

Do geometric invariants of program graphs add information over classic semantic analysis?

MANIFOLD: a falsifiable study on OWASP Benchmark 1.2

MANIFOLD Project

Working draft v0.4 — not peer reviewed

Abstract

MANIFOLD is a *neuro-symbolic structural-invariant verifier*: an autonomous agent that mines security specifications (authorization invariants, state-machine cycles, taint preconditions) from code with an LLM and discharges them with formal methods (Z3, persistent/directed homology, a mechanized Lean 4 core) — emitting only *machine-checkable witnesses under the modeled semantics*. This paper asks one integral question and answers it with a sharp boundary: *where do the structural invariants of program graphs carry security information, and where do they not?*

The negative side of the boundary is precise. On linear injection flaws, where a flat, near-acyclic servlet has no curvature or persistent topology to exploit and classic taint is already near-optimal, structural priors add nothing: on OWASP Benchmark 1.2 (2740 cases) raw taint ranks the vulnerable sink best (MRR 0.894, AUC 0.895, beats the fused field with $p = 5 \times 10^{-5}$, add-one over 20000 permutations); spectral (0.496) and topological (0.500) signals are at chance and curvature below chance (0.376). Hypotheses H1/H2/H3/H5 are *not supported*. The positive side is where the highest-impact flaws of modern software live — BOLA/IDOR (CWE-639), missing authorization (CWE-862), and destructive state cycles (reentrancy, deadlocks) — which are *structural*: the offending data is legitimate, so classic taint is blind to them, and only the privilege-boundary and state-transition *shape* of the program gives them away. There the same machinery is the only signal: on a curated 39-case BOLA/IDOR corpus the structural *missing-authorization* detector reaches recall 1.0 / precision 0.72 / F1 0.837 while taint scores 0 (its 7 false positives are the documented over-approximation limit); and GLMY directed path homology extracts the concrete β_1 generator of a state cycle that symmetrization erases. The result is a single account — the same typed graph, privilege algebra, directed homology and Z3 verifier — of *which* program relations make structural invariants informative. We contribute this account, a prioritization-oracle harness, a reproducible Java pipeline with Z3 witnesses, external validation on Juliet and real CVEs, and a Lean 4 mechanization of the taint-abstraction soundness and the no-auth boundedness part of the authorization invariant.

1 Motivation

Rule-based and semantic static-analysis pipelines are brittle: they miss novel vulnerabilities and drown analysts in false positives. Two capabilities are missing:

- (1) A *unified, quantitative representation* of “program structure” onto which many *independent* signals can be projected and fused.

- (2) A *reasoning loop* that converts a “suspicious region” into a “verified finding under the encoded semantics” — the LLM supplies the intuition, mathematics supplies the verification.

MANIFOLD provides both.

2 The core idea




Figure 1: The pipeline: an artifact becomes a typed graph, enriched with mathematical structure into a vulnerability scalar field $V(x)$, which an LLM agent navigates; a formal verifier confirms or refutes each hypothesis and the result re-embeds into the manifold.

We model vulnerability *candidates* as violations of program invariants rather than purely syntactic patterns (a broken authorization invariant, a persistent cycle, a negative curvature chokepoint, a taint flow across a spectral trust cut). Each layer detects a different *kind* of invariant violation, and the fused field $V(x)$ ranks regions for the agent to inspect.

3 The Intermediate Representation (IR)

Definition 1 (Program graph). A *program graph* is a tuple $G = (V, E, \tau_V, \tau_E, \omega)$ where

- V is a finite set of *nodes* (modules, functions, blocks, statements, calls, variables, sources, sinks);
- $E \subseteq V \times V$ is a set of directed *edges*;
- $\tau_V : V \rightarrow \mathbb{N}_K$ assigns each node a *kind* from $K = \{\text{module, function, block, statement, call, assign, variable, parameter, source, sink, gate}\}$;
- $\tau_E : E \rightarrow \mathcal{T}$ assigns each edge a *type* $\mathcal{T} = \{\text{control, data, call, taint, trust, auth}\}$;
- $\omega : E \rightarrow \mathbb{R}_{\geq 0}$ assigns each edge a non-negative *weight* (default 1).

The IR is the single source of truth. Every adapter (source code, binary, web API, LLM agent) must produce a program graph; every analysis consumes it. The schema is language-agnostic and JSON-serializable (the Python and Rust layers share the identical schema).

Naming note. Despite the codename, the object is a *program-analysis space*: a typed graph together with derived complexes, embeddings and a scalar field. We do *not* claim it is a mathematical manifold (there is no atlas, chart or local Euclidean structure); “manifold” is used informally, as in “data manifold”.

4 Mathematical layers

For each layer we fix a subgraph by edge type τ_E , then compute structure.

4.1 Spectral layer (algebraic graph theory)

For an edge type t , define the symmetric *adjacency* A_t and *degree* D_t , and the *combinatorial Laplacian*

$$L_t = D_t - A_t.$$

L_t is positive semidefinite; its spectrum $0 = \lambda_1 \leq \lambda_2 \leq \dots$ encodes global structure:

- λ_2 is the *algebraic connectivity* (Fiedler value). Its eigenvector f_2 — the *Fiedler vector* — induces a canonical cut $S^+ = \{v : f_2(v) \geq 0\}$.
- The *spectral embedding* $\Phi^{(d)}(v) = (f_2(v), f_3(v), \dots, f_{d+1}(v))$ maps nodes into $\mathbb{R}^d$.
- The Cheeger inequality bounds the bottleneck ratio $h(G)$ by $\frac{\lambda_2}{2} \leq h(G) \leq \sqrt{2 d_{\max} \lambda_2}$.

Signal 1 (spectral anomaly). The reconstruction error of a node under a low-rank eigenmap reconstruction, and its distance from the Fiedler cut, are used as anomaly scores.

4.2 Topological layer (persistent homology)

From the program graph we build a *filtration* of simplicial complexes (e.g. a Vietoris–Rips or flag complex over the shortest-path metric, or a filtration driven by ω). Persistent homology yields, for each dimension k , a *persistence diagram* Dgm_k of birth–death pairs. β_0 counts components, β_1 counts independent cycles; points far from the diagonal have high persistence. The **Mapper** algorithm (Singh–Mémoli–Carlsson) builds a navigable *nerve* from a filter function $f : V \rightarrow \mathbb{R}$, producing the map that gives the project its name.

4.3 Geometric layer (discrete curvature)

The *Ollivier–Ricci curvature* of an edge (u, v) is

$$\kappa(u, v) = 1 - \frac{W_1(m_u, m_v)}{d(u, v)},$$

with lazy random-walk measures m_u, m_v and Wasserstein-1 distance W_1 . Negative curvature marks *funnels* through which many shortest paths pass. The *Forman–Ricci* curvature (unweighted) is $\kappa_F(u, v) = 4 - \deg(u) - \deg(v)$, cheaply flagging high-degree hubs.

4.4 Algebraic layer (taint lattice & authorization order)

Taint lattice. The abstract domain is the powerset lattice of taint tags $\mathbb{T} = (2^{\mathcal{K}}, \subseteq, \cup, \cap)$, with $\perp = \emptyset$ (clean) and $\top$ (tainted). It forms a Galois connection (α, γ) with the concrete collecting semantics; soundness is proved in §4.6. *Authorization order.* We keep the authorization model deliberately minimal — no category theory is needed. Let $F : V \rightarrow \mathbb{N}$ be a *privilege-level function* on program points (so different points may share a level; an edge need not raise it); an edge $e = (u, v)$ is an *auth edge* iff $F(u) < F(v)$ (a strict escalation), and E_A is the set of auth edges. A *sanitizer* at a point maps every tag to clean. **Authorization invariant:** a taint flow reaching a sink is a *violation* when its path crosses an auth edge and no sanitizer on the crossing segment restores the invariant. This is stated precisely in §4.6 (the no-auth boundedness part is mechanized).

4.5 Formal layer (symbolic execution & model checking)

Symbolic execution computes a path condition Φ and symbolic state σ ; an SMT solver checks $\text{SAT}(\phi_{\text{bad}})$, whose model is a witness under the encoded semantics (not an end-to-end exploit). Reachability is the CTL formula $\mathbf{EF} \phi_{\text{bad}}$; taint propagation is a finite automaton whose accepted language is exactly the set of source $\rightarrow$ sink paths.

4.6 Formal core: authorization invariant and soundness

Small-step semantics of the IR. A state maps variables to concrete values; the transition relation $v \xrightarrow{\sigma} v'$ steps through the IR. The concrete taint collecting semantics $\mathcal{C}$ maps each point to the set of executions that carry an untrusted value there. The abstract analysis of §4 is the function A computed by the transfer functions; the pair (α, γ) satisfies $\alpha(\mathcal{C}(v)) \sqsubseteq A(v)$ and $\mathcal{C}(v) \subseteq \gamma(A(v))$.

Proposition 1 (soundness of the taint abstraction). *If the abstract analysis reports a variable clean, no concrete execution taints it.*

Proof scheme. By induction on the transition sequence, using the invariant “concrete taint $\subseteq$ abstract taint” at every state. The base holds for entry variables; the step holds because each transfer function is monotone and over-approximates its concrete counterpart (assignments join operand tags, sanitizers map to clean, branch joins upper-bound the merge). The guarantee is contingent on the sanitizer *summary* matching the sanitizer’s real semantics: a wrong summary (claiming to clean what it does not) breaks soundness. This is exactly the mechanized statement `prop1_soundness` (§4.6), proved for a concrete boolean abstraction; extending it to the full tag lattice is straightforward but assumed here.

Proposition 2 (authorization-invariant violation). *Let $F : V \rightarrow \mathbb{N}$ be the privilege-level function, and let $p = e_1 \circ \dots \circ e_n : s \rightarrow k$ be a path. Write $\text{Auth}(p)$ for the auth edges on p and $\text{San}(p)$ for the sanitizers on p . Then: (i) if $\text{Auth}(p) = \emptyset$ then $F(k) \leq F(s)$ (no privilege escalation); (ii) if an auth edge $e_i = (v_i, v_{i+1})$ lies on p and no sanitizer neutralizes the tag before that crossing, then the execution contains an attacker-controlled privilege-boundary crossing.*

Proof. (i) Each non-auth edge satisfies $\neg(F(u) < F(v))$, i.e. $F(v) \leq F(u)$, so F is non-increasing along p ; transitivity gives $F(k) \leq F(s)$. This is the mechanized `no_auth_bounded` (§4.6). (ii) An auth edge strictly raises F at the crossing; absent a sanitizer before it, the tag survives, so the boundary is crossed under attacker control. The statement is *local* to the crossing: mere presence of an auth edge on p does not imply $F(k) > F(s)$ — e.g. $0 \rightarrow 1 \rightarrow 0$ crosses an auth edge yet returns to the same level. Bounding $F(k)$ (or placing k inside the protected region) needs the additional assumption $F(k) > F(s)$. (A formula of the form $F(p) \circ \eta_s \neq \eta_k \circ T(p)$ appeared in an earlier draft; it was under-typed and is withdrawn — no natural transformation is claimed here.) $\square$

The invariant (ii) is exactly the operational rule H4 (§8/C4): a taint finding inside an authenticated function is an authorization violation. Proposition 1 and the no-auth boundedness part (i) of Proposition 2 are *mechanized* in Lean 4 (core only, no Mathlib) in `formal/Manifold.lean`: `prop1_soundness` proves the boolean taint abstraction sound against the tag-collecting semantics (§3 there), and `no_auth_bounded` proves (i) (§4 there). The crossing part (ii) is a semantic definition, not separately mechanized. The proof checking is part of the reproducibility artifact (`scripts/check_formal.sh`).

5 The vulnerability scalar field

$$V(x) = \sigma\left(\sum_i \alpha_i s_i(x)\right),$$




Figure 2: The manifold: nodes are program components, edges are typed relations. A taint flow (red) into a sink raises $V(x)$, highlighting a region for the agent to inspect and the verifier to confirm.

where s_i are normalized layer signals and α_i weights. To avoid a *circularity* — the field that decides *what to verify* must not depend on verification — we distinguish two fields:

- $V_{\text{prior}}(x)$ fuses only signals computable *before* SMT (**taint**, spectral, topological, geometric). This is the field used for ranking and for the prioritization experiment (§12.1).
- $V_{\text{post}}(x)$ additionally consumes the **formal** (Z3) signal and is used only for re-embedding / post-hoc analysis; it never feeds the pre-verification ranking.

In the prototype the weights are *uniform* ($\alpha_i = 1/N$ over the chosen signals) and V is reported as the raw fused score; calibrating α_i is future work.

6 Mapping hypotheses

#	Law	Vulnerability class
H1	Chokepoint ($\kappa \ll 0$)	privilege escalation, auth bypass
H2	Cycle (persistent H_1)	reentrancy, infinite recursion
H3	Trust cut (Fiedler cut)	injection across trust boundary
H4	Authorization-invariant violation	arbitrary taint violation
H5	Persistence outlier	real vs. spurious features
H6	Reachability ($\text{SAT}(\phi_{\text{bad}})$)	model witness (encoded semantics)

7 A worked example (end-to-end)

Consider this ten-line Python fragment:

```
def search(db, query_param):
    sql = "SELECT * FROM users WHERE name = '" + query_param + "'"
    cursor = db.cursor()
    cursor.execute(sql)
```

The pipeline produces: (i) an IR with a **source** node (`param:query_param`), a **sink** node (`cursor.execute`, category `sql`) and a **taint** edge between them; (ii) the field $V(x)$ spikes at the sink (taint + formal signals active); (iii) the Z3 symbolic verifier returns **SAT** with a model witness under the encoded semantics (e.g. `"x' OR '1'='1"`), proving the sink's argument depends on the

attacker-controlled symbol. The parameterized variant `execute(query, (param,))` is correctly *not* flagged, and a `@login_required` gate on the enclosing function upgrades the finding to an H4 authorization violation (a taint flow crossing an `auth` edge).

8 The autonomous LLM agent

A closed loop: (1) *ingest* into the IR; (2) *embed & map* ($V(x)$, spectral embedding, persistence, curvature, Mapper); (3) *hypothesize* (the LLM proposes candidates from the manifold); (4) *verify* (the formal layer proves or refutes); (5) *refine* (findings re-weight and re-embed); (6) *report*. The LLM provider layer targets any OpenAI-compatible endpoint via **LiteLLM**. The LLM does not do raw detection: mathematics proposes regions, the LLM proposes hypotheses, the formal layer proposes proofs.

9 Design challenges and mitigations

A design review identified five structural risks. For each we state the mitigation adopted in the prototype and the open problem.

C1 – The bottleneck paradox. Negative curvature flags funnels, but healthy software *deliberately* concentrates flow through `auth` middleware, sanitizers, dispatchers and loggers. Reporting these as privilege-escalation leads yields systematic false positives. *Mitigation:* the agent discriminates defensive chokepoints (by role) and drops or down-ranks them; the principled fix is to tag `trust/auth` edges (already reserved in the IR) and treat a chokepoint as defensive iff it is the target of an `auth` edge. *Open:* learning the defensive/offensive label from data.

C2 – Symmetrization destroys causality. The combinatorial Laplacian $L = D - A$ and standard persistent homology act on undirected metric spaces, but control and data flow are directed: A tainting B is not B tainting A . Symmetrizing can destroy the direction of taint. *Mitigation:* the prototype implements Chung’s directed Laplacian restricted to the largest strongly connected component; the Perron vector π is computed by power iteration with a small PageRank-style teleportation (damping $\alpha = 0.01$), which keeps the stationary distribution well-posed on periodic chains. This operator requires *recurrent* (strongly connected) structure; on a DAG it yields no informative signal. A DAG certainly has direction — its causality is simply delegated to the directed taint/reachability analysis (and, for homology, to GLMY path homology, §11) rather than to a stationary-distribution operator, which cannot represent it without additional structure. *Implemented:* GLMY directed path homology H_0, H_1 (Grigor’yan–Lin–Muranov–Yau), which changes the verdict versus the symmetrized homology (§11).

C3 – Cost of optimal transport. Ollivier–Ricci requires a transportation LP per edge: $\mathcal{O}(|E| \cdot \text{LP})$, prohibitive beyond $\sim 10^5$ edges. *Mitigation:* Forman–Ricci $\kappa_F = 4 - \deg(u) - \deg(v)$ is $\mathcal{O}(|E|)$ and is the default chokepoint signal; entropy-regularized (Sinkhorn) transport is implemented as the fast alternative to the exact LP. *Open:* adaptive regularization schedules for very large graphs.

C4 – Operationalizing H4 (authorization invariant). A constructive rule is required. Recall the privilege-level function $F : V \rightarrow \mathbb{N}$ and the `auth` edges $E_A = \{(u, v) : F(u) < F(v)\}$ from §4.4; the extractor only needs to emit the `auth` edges that mark privilege boundaries. Then a taint flow $s \rightarrow^* k$ *violates the authorization invariant* iff its path crosses at least one `auth` edge — i.e. untrusted data enters a protected region without the boundary neutralizing it. This is a reachability query over the combined (data $\cup$ taint $\cup$ `auth`) graph, directly computable and automatable. *Implemented:* an authorization extractor emits `gate` nodes and `auth` edges from

privilege-gate decorators, and a taint finding inside a gated function is reported as an authorization violation.

C5 – Manifold-to-verifier bridge. Geometric features are quantitative; Z3 needs exact preconditions. We fix a *grammar*: a hypothesis is the structured spec (`sink`, `source`, `line`, `category`). The manifold supplies the spec (it knows the sink and line), the LLM supplies the interpretation (CWE, description, confidence), and the verifier consumes the spec — never the LLM prose. This keeps the SMT query precise and reproducible.

The thesis our data supports is more modest. Semantic dataflow (taint) is the dominant signal on the evaluated corpora; structural invariants provide *complementary hypotheses*. The open question is whether structural priors can reduce *verification cost* in regimes where semantic analysis is incomplete (interprocedural, framework-mediated), and whether they help at very small budgets.

10 Implementation status

Phase	Deliverable	Status
F0	Whitepaper, architecture	done
F1	IR + Python SAST ingest + taint	done
F2	Spectral + geometric kernels (Rust)	done
F3	TDA (persistent homology, Mapper)	done
F4	Formal verification (taint lattice, Z3)	done
F5	Autonomous LLM agent (LiteLLM)	done
F6	Visualization + end-to-end demo	done
F7	Binary / web / LLM adapters	done

11 Threats to validity

Internal. The Java adapter is intraprocedural; OWASP’s distractor flows and Juliet’s cross-class variants are therefore measured with an inherent recall gap. Ground truth is the benchmark’s own (OWASP’s “false” cases are partly definitional and debatable). *External.* Results cover Java servlets and the Juliet suite; Python, binaries and LLM agents are implemented but not benchmarked at the same scale. *Construct.* “Vulnerable” is a per-case label; the field is per-node, so the prioritization experiment is a proxy, not a downstream exploitability measure. *Conclusion.* The negative result (geometric priors do not beat taint) is specific to these corpora and signals; it should be re-tested on path-sensitive, interprocedural analyses and on real CVEs before being generalized.

12 Validation and responsible use

Validation is two-stage. Stage 1 (complete) runs the toy micro-benchmarks of F1–F7 against known ground truth: every seeded vulnerability is confirmed and every sanitizer / parameterized-query / argv-form false positive is excluded. Stage 2 is complete for the Java/OWASP evaluation and ongoing for SARD, binaries and additional domains; it measures precision/recall of the fused field versus raw taint over labeled corpora — experimental validation at scale, distinct from the completed prototype. The harness is implemented: a labeled toy benchmark reports precision/recall/F1 for the raw taint pass, the Z3-verified pass and a *calibrated* field whose weights α_i are learned by leave-one-out logistic regression over the cheap signals — on the toy corpus the calibrated field recovers the reentrancy cycle (topological signal) that the taint pass misses, raising recall while

keeping precision at 1.0. A corpus loader accepts a JSON manifest or a “vulnerable|safe“ directory, so the harness can be pointed at OWASP Benchmark / SARD exports; and the binary adapter has an optional angr backend (“CFGFast“ + interprocedural taint) in addition to the zero-dependency objdump path. Running the Java SAST adapter over the full OWASP Benchmark 1.2 (2740 test cases) gives the naive baseline: precision 0.515 / recall 0.426 / F1 0.466 on the seven taint categories, and precision 0.544 / recall 0.634 / F1 0.586 over all eleven categories — the reference point that the fused field, calibration and the verifier are meant to improve on. MANIFOLD is a defensive / authorized-testing tool; we release under the assumption of authorized use and encourage coordinated disclosure.

12.1 The prioritization oracle (main experiment)

The bottleneck of SAST is not *detecting* sinks but *deciding which one to verify*. For each of the 902 vulnerable taint cases we enumerate *every* sink call (the candidate set), rank the candidates, and measure the rank at which a ground-truth-vulnerable sink appears. Baselines: DFS (source order), raw taint, random (20 seeds, 95% CI), and the fused field $V(x)$.

Ranking	MRR	R@1	mean rank	NDCG@5
DFS (source order)	0.848	0.755	1.517	0.882
raw taint	0.894	0.847	1.422	0.916
$V(x)$ (fused field)	0.872	0.805	1.475	0.899
random	0.709 ± 0.042	0.531	—	0.778

The honest, decisive finding. The fused field does *not* beat raw taint: it is slightly *worse* (0.872 vs. 0.894 MRR). The per-signal AUC for predicting the vulnerable sink shows why:

Signal	AUC [95% CI]
taint	0.895 [0.881, 0.909]
formal (Z3/oracle)	0.878 [0.865, 0.891]
spectral (Fiedler)	0.496 [0.469, 0.518]
topological (H_1)	0.500 [0.500, 0.500]
geometric (Forman κ)	0.376 [0.360, 0.394]

The spectral and topological signals are *indistinguishable from chance*, and the curvature signal is *worse than chance* (negative curvature flags high-degree hubs, which are *not* the target sinks). Per the mapping-hypotheses program (ġ4/ġ6), this is a publishable *negative* result: on OWASP Benchmark, **the geometric and topological priors do not survive beyond taint**. The structural prior does not improve average prioritization; MANIFOLD’s remaining value lies in formal witnesses, the unified experimental substrate, and the diagnosis of semantic-analysis bottlenecks. H1, H2, H3 and H5 are therefore reported as *not supported* on this corpus. For H3 this rests on the spectral (Fiedler) signal being at chance (0.496); note that the evaluated spectral score aggregates low-rank reconstruction error with Fiedler-cut distance, so it is a proxy for the trust-cut predicate rather than an isolated test of it.

Paired test and global pool. The difference taint – field is +0.0225 (MRR), 95% bootstrap CI [+0.0157, +0.0294], permutation $p = 5 \times 10^{-5}$ (add-one estimator over 20000 permutations, 2000 bootstrap): taint significantly dominates. To also measure the real analyst situation (safe and vulnerable candidates competing in one queue), we pooled *every* sink candidate over all 1698 taint-category cases and ranked globally:

Global ranking	AP (PR-AUC)	P@10	R@10	queries to 90% recall
DFS (source order)	0.286	0.2	0.002	3596
raw taint	0.489	0.6	0.006	2080
V_{prior} (structural)	0.434	0.8	0.008	3207
random	0.216	0.2	0.002	4176

The picture is nuanced: the structural field is best at the extreme top ($P@10 = 0.8$ vs. 0.6) but has lower average precision and needs more queries to reach 90% recall. The $P@10$ gap is small in absolute terms — eight versus six correct items in the top ten — so we do not read it as a robust advantage; it is consistent with tie-breaking noise and we report it without a significance claim. The overall verdict remains negative: the geometric/topological priors help only marginally at very small budgets and hurt on average — a sharper, still negative, version of the H1/H2/H3/H5 verdict.




Figure 3: Cost curve: cumulative recall vs. verification budget

Signal	AUC [95% CI]
taint	0.895 [0.881, 0.909]
formal	0.878 [0.865, 0.891]
spectral	0.496 [0.469, 0.518]
topological	0.500 [0.500, 0.500]
geometric	0.376 [0.360, 0.394]

12.2 Empirical study on OWASP Benchmark 1.2

We ran the Java SAST adapter over the full corpus and profiled the false negatives. Three *generalizable* gaps explained most of the loss, and were fixed (none is tuned to the benchmark):

- (i) **Branch-aware taint.** A sequential, path-insensitive pass let a null guard (`if (param == null) param = "";`) clear the taint, even though the tainted value survives on the other path. We now snapshot both branches and *join* their environments (union of tags), the taint analogue of a lattice least-upper-bound.
- (ii) **Receiver and state taint.** Sinks such as `statement.execute()` and `pb.start()` take no argument — the taint lives in the *receiver* object, which was obtained from a tainted constructor/mutator. We propagate taint from the receiver, and mark a variable tainted when a mutator (`add/put/command/...`) stores a tainted value into it.

- (iii) **Sink coverage.** The XSS family includes `printf/format`, and only writes to the *HTTP response writer* are XSS — `System.out.println` of tainted data is not. Restricting the XSS sinks to the response writer removed a large class of false positives.

The ablation isolates each contribution (precision P , recall R , F1; taint categories, 1698 cases):

variant	P	R	F1
initial naive adapter	0.515	0.426	0.466
+ branch-join, receiver/state taint, sink coverage	0.530	0.856	0.655
+ XSS response-writer restriction (final)	0.549	0.834	0.662
final, without branch-join	0.515	0.555	0.535
final, without XSS restriction	0.530	0.856	0.655

Branch-aware taint is the dominant recall lever (+0.28 recall); the response-writer restriction is the precision lever. The final adapter reaches $P = 0.549$, $R = 0.834$, $F1 = 0.662$ on the taint categories and $P = 0.572$, $R = 0.842$, $F1 = 0.681$ over all eleven (crypto/hash are configuration-driven and partly out of scope). The *residual* errors are the benchmark’s deliberately adversarial “false” cases — debatable parameterized-query and `ProcessBuilder`-list cases — and cross-category distractor flows: precisely the precision that the topological/spectral signals, the Z3 verifier and the calibrated field are designed to recover.

12.3 Comparison with industrial SAST

The official OWASP Benchmark score is the *Youden index* $J = \text{TPR} - \text{FPR}$ (the *Benchmark Accuracy Score*, i.e. TPR minus FPR). Table 1 places MANIFOLD against two widely deployed tools. SonarQube and CodeQL rows are *reported* numbers from public OWASP Benchmark 1.2 evaluations; both MANIFOLD rows are our *measured* results. **Caveat:** the external figures are not re-executed here and may use different aggregation conventions (some are macro-averages per CWE, others micro or per-rule), so the mixed precision/F1 columns are *not* directly comparable across rows; only the TPR/FPR ordering should be read, and a uniform re-execution of CodeQL/-SonarQube under one protocol is required before any head-to-head claim.

Tool / config.	TPR	FPR	Prec.	F1	J	Time
SonarQube (SAST) (a)	0.956	0.946	0.330	0.490	+0.010	3 min
GitHub CodeQL (default) (a)	0.902	0.682	0.603	0.744	+0.220	18 min
MANIFOLD adapter (measured, 11 cat.) (b)	0.842	0.674	0.572	0.681	+0.168	45 s
MANIFOLD adapter + Z3 (measured, 7 taint cat.) (b)	0.834	0.776	0.549	0.662	+0.057	4 min

Table 1: OWASP Benchmark 1.2 (2740 cases: 1415 vulnerable / 1325 safe). (a) reported in public evaluations, not re-run here. (b) **measured**. The Z3 verifier reproduces the adapter’s precision (it adds per-finding *witnesses*, not precision): the residual errors are adversarial/definitional cases, not spurious symbolic paths. The $V(x)$ prioritization prior is evaluated separately in §11.1 (it does not beat raw taint).

Methodology notes. Hardware: 8-core x86-64, 32 GB RAM, NVMe SSD. MANIFOLD: Python 3.14, `tree-sitter-java`; Z3 5.x with a *5 s per-query timeout* — on timeout the case is marked *unknown/conservative*, never counted as a detection, so the projected precision is not inflated by aborts. CodeQL: `java-security-extended.qls` on a clean database. SonarQube: Community

Edition with SonarJava security rules. CWE mapping: the eleven categories of Benchmark 1.2; `crypto/hash` (CWE-327/328) are configuration-driven and require a static-configuration extractor into the IR. The measured adapter flags a case vulnerable iff it emits any finding; precision is bounded by the benchmark’s adversarial “false” cases (debatable parameterized-query and `ProcessBuilder`-list patterns) and by cross-category distractor flows — the gap the neuro-symbolic rows target.

12.4 Closing the Z3 row (Java verification)

We implemented a Z3-backed Java symbolic verifier (bounded, branch-merging, string semantics; sanitizers as contracts) that confirms or refutes each taint flow and emits a machine-checkable *witness* under the modeled semantics. Measured on the seven taint categories (1698 cases) it reproduces the adapter exactly ($P = 0.549$, $R = 0.834$, $F1 = 0.662$, $J = +0.057$) while adding per-finding witnesses. The honest conclusion: **the formal layer adds machine-checkable witnesses under the *modeled* semantics, not precision, on this benchmark** — a witness proves satisfiability inside the model (which abstracts library summaries and framework behaviour), not end-to-end exploitability of the real program. The residual false positives are adversarial/definitional (debatable parameterized-query cases), not spurious symbolic paths; catching them requires modelling the benchmark’s own semantics, not stronger SMT.

12.5 Cross-corpus transfer (Juliet)

To test transferability we ran the OWASP-tuned adapter on the *Juliet Java* test suite (find-secbugs mirror; each test case concatenated with its flow-variant files). On 264 cases across five CWEs it scores $P = 0.333$, $R = 0.562$, $F1 = 0.419$: detection is partial and precision collapses on the “good” variants. The reason is structural, not a tuning artifact — Juliet splits source and sink across `_a/_base/_bad` classes, so correct analysis requires *interprocedural* taint, which the intraprocedural adapter does not provide.

12.6 Interprocedural taint: necessary, not sufficient

We added a context-insensitive interprocedural fixpoint (`java_interproc.py`): a callee parameter becomes tainted if some call site passes tainted data to it, and a call becomes tainted if its callee returns tainted data; the fixpoint is iterated to convergence. On Juliet it raises recall $0.562 \rightarrow 0.750$ and $F1\ 0.419 \rightarrow 0.490$ (CWE78/CWE90 recover from 0.000 to 0.750). *On OWASP it is neutral* (taint $F1\ 0.662 \rightarrow 0.656$): the interprocedural pass adds recall ($0.834 \rightarrow 0.856$) but also coarse parameter-driven false positives, so precision drops, and we union the pass with the full intraprocedural adapter for coverage parity. The lesson is precise: interprocedural reachability is *necessary* for cross-class flows but *not sufficient*. We then added *caller-context sensitivity* with per-context return summaries: it provably prunes a synthetic class of return-driven false positives (a value returned from a tainted context no longer taints a clean caller’s call), yet it leaves OWASP and Juliet *unchanged* ($F1\ 0.656$ and 0.490). The real levers are therefore not context but (a) framework/library summaries and (b) fine-grained return precision — not more geometry.

12.7 Framework summaries

Guided by the CVE evidence, we added an explicit framework/library summary layer: *attribute-style sources* (`request.GET`, `request.POST`, `request.args`, `request.form`, `self.request.POST`) and *ORM/template sinks* (`RawSQL`, `raw`, `extra`, `text`, `render_template_string`, `JPA`

`createQuery`). With *positional* sink summaries (only `arg 0` of a query/code call is dangerous), the analyzer now sees application-level flows it previously missed — Django query-parameter SQL injection, Flask SSTI, JPA injection — while keeping parameterized queries and literal templates safe. This is the concrete lever the CVEs pointed to. It does not move the aggregates, but for two different reasons: OWASP/Juliet failures are mostly *structural* (cross-method flows that need interprocedural analysis), whereas the unresolved CVEs are *framework-internal* (the vulnerable logic lives inside the framework itself, with no application-level source).

12.8 Scalability of each kernel

We timed every kernel on synthetic graphs (log-log, Fig. 4). **Forman–Ricci** (0.6 ms @ 10^3 edges $\rightarrow$ 55 ms @ 10^5) and **Mapper** (0.5 $\rightarrow$ 56 ms) are near-linear. **Persistent homology** is super-linear (11 ms @ $10^3 \rightarrow$ 1.1 s @ $10^4 \rightarrow$ 123 s @ 10^5) because the fundamental-cycle extraction BFSes the spanning forest per cycle; a persistence-aware union-find would make it near-linear. **Sinkhorn** (0.71 s @ $10^3 \rightarrow$ 16.4 s @ 5×10^3) and **exact Ollivier–Ricci** (34 ms @ 200 $\rightarrow$ 720 ms @ 10^3) are quadratic, so the latter is only for small graphs (the reason Forman–Ricci is the default, §8/C3). The **dense spectral and directed Laplacians** are cubic (26 ms @ 300 $\rightarrow$ 1.7 s @ 1200 nodes) and *cannot* reach 10^5 nodes without the sparse solvers planned in F2. The practical envelope today is $\sim 10^5$ edges for the near-linear signals, which covers the corpora used here.




Figure 4: Scalability of each kernel (log-log)

12.9 Real CVE fixes

We ran the vulnerable-vs-patched comparison on real CVE fix commits drawn from OSV / GitHub Security Advisories (`benchmarks/run_cves.py`; the harness reconstructs the *before* revision by reverse-applying the commit diff to the commit’s raw content, so it needs no API quota). On 22 CVEs and 46 changed source files, **3 files resolved a finding**: they are `eval` injections — `senaite.core` (CVE-2026-54569; `value = eval(value)` on a method parameter) and `xinference` (CVE-2026-61539; `eval(text)` on LLM output) — plus a partial resolution in `nltk` (CVE-2026-79676) — flagged in the vulnerable revision and clean in the patch. The remaining CVEs are mostly *framework-internal* or non-taint fixes (Django ORM SQL injection, Jinja, requests, GitPython, NLTK argument-injection and integrity checks): the flow is indirect or outside our sink model, so the adapter sees nothing. This mirrors the Juliet result and locates the frontier precisely: MANIFOLD

already catches application-level `eval/exec/deserialization` fixes; framework-mediated flows need interprocedural, context-sensitive, library-aware analysis.

12.10 Directed path homology (GLMY)

We implemented Grigor’yan–Lin–Muranov–Yau *path homology* H_0, H_1 over the directed 1-skeleton (`core/src/path_homology.rs`): Ω_2^∂ is the set of allowed 2-paths whose long edge is allowed and non-degenerate ($u \neq w$), and ranks are computed by Gaussian elimination. Direction changes the verdict. On `reentrancy.py`, undirected homology reports $\beta_1 = 1$ (the `withdraw/read_balance/update_balance` triangle), while GLMY reports $\beta_1 = 0$: the transitive 2-paths through the apex `withdraw` fill the `read_balance`↔`update_balance` 2-cycle. On an isolated 2-cycle $A \leftrightarrow B$ the roles invert (GLMY $\beta_1 = 1$, undirected $\beta_1 = 0$), and the analyzer returns the H_1 *generator* itself — the directed cycle as the edge list $\{A \rightarrow B, B \rightarrow A\}$ — so the invariant is actionable, not just a number. Symmetrization (C2) is therefore *not* innocuous: the directed invariant is a genuinely different, and sometimes opposite, signal.

12.11 LLM ablation

We ran the agent twice over a varied corpus of 38 Python files (29 with a confirmed finding, spanning command / code-exec / SQL / deserialization / path-traversal / file-write / logging / SSTI, plus safe and gated variants): offline (heuristic hypothesizer) and online through the local **LiteLLM** proxy (DeepSeek `flash`, OpenAI-compatible /v1). The *detection* and *ranking* axes are invariant: the confirmed finding sets and their confidence ordering agree on 37/38 files (agreement 0.974), because the SMT verifier — not the LLM — is the arbiter. The single mismatch is a `compile` code-exec case where the LLM’s hypothesis is itself nondeterministic (sometimes empty), so it occasionally fails to re-propose the finding rather than changing it. The *interpretation* axis changes: on 8 of the 29 confirmed findings (28%) the LLM assigns a different, usually more specific CWE (e.g. CWE-674 “Uncontrolled Recursion” for the reentrancy cycle vs. the heuristic’s generic label), and produces more concise titles. This larger controlled ablation (38 files, `benchmarks/run_llm_ablation.py`) confirms the arbitration boundary at scale: *the LLM adds interpretability and reporting, not detection or ranking* — the neuro-symbolic split holds with a small, measured nondeterminism leak. A Java/OWASP-side ablation is future work.

12.12 Dual binary backend

The binary adapter ships two interchangeable backends that emit the same IR: (i) a zero-dependency `objdump`-text path (function labels + `call` edges + interprocedural taint reachability) and (ii) an *optional* `angr` backend (CFGFast recovery of the interprocedural call graph + the same taint reachability), selected by `-adapter angr-binary`. The fast path keeps MANIFOLD usable without heavy dependencies; the `angr` path recovers control flow that disassembly heuristics miss. This parity is the substrate for future ELF/PE benchmarks from SARD.

13 Structural invariants across domains

The two results so far are two sides of one boundary, not two projects. On one side (linear injections) a flat, near-acyclic servlet has no curvature or persistent topology to exploit, and classic taint is already near-optimal — so structural priors add nothing. On the other side the offending data is legitimate and only the *structure* of the program — its privilege boundaries and state transitions

— reveals the flaw, so structure is the *only* signal. The same machinery (typed graph, privilege algebra, directed homology, Z3) spans both; what changes is the program relation being modelled. This section surveys the four domains where conventional static analysis is blind and where each mathematical layer has a natural home.

13.1 Authorization invariants: BOLA / IDOR / BFLA

Broken Object Level Authorization (BOLA/IDOR, CWE-639) and missing authorization (CWE-862) top the OWASP API Security Top 10, and classic taint cannot see them: the object selector is a *clean* value (an integer id, often cast), the sink is an ordinary query, and the bug is that user A reads user B’s data — there is no “dirty” data to track. The modeled object is a *multipartite* graph: the API surface (OpenAPI operations), the backend endpoints, the permission decorators (*gate* nodes) and the data models, as produced by the OpenAPI adapter (`manifold.ingest.web`), which already emits `auth` edges from security schemes. The signal is the complement of the mechanized `no_auth_bounded` invariant: a directed path from a public entry point to a protected resource with *no* strict authorization edge where one is required. MANIFOLD implements this as the *missing-authorization / BOLA* detector (`manifold.analysis.authorization.bola_idor_candidates`), and Z3 then proves whether a session assignment (`session.user_id  $\neq$  resource.owner_id`) reaches the resource — a *mathematical BOLA witness*. On a curated corpus of 39 cases (`examples/python/bola/`, `examples/python/bola_corpus/`, evaluated by `benchmarks/run_bola.py`) taint reports *zero* findings on the vulnerable cases while the structural detector reaches **recall** 1.0 at **precision** 0.72 (F1 0.837): a feasibility signal with an honest precision limit, because the structural heuristic over-approximates (parameterized/sanitized/unused-source cases are the 7 documented false positives). It is precisely the regime where taint provably scores 0 and structure is the only signal.

13.2 State machines, concurrency and protocols: directed homology

In a servlet there are no state cycles; in smart contracts, lock-acquisition protocols and concurrency controllers the worst flaws *are* destructive feedback loops (reentrancy, deadlocks, illegal transitions such as `Unauthenticated` $\rightarrow$ `Authenticated` skipping signature validation). The modeled object is the *state-transition graph* (the interprocedural CFG extended with external calls and global-store mutations). The persistent directed H_1 generator of GLMY path homology is the exact signature of a reentrancy cycle — and, as §4.2 showed, symmetrization is *not* innocuous: undirected homology can report $\beta_1 = 1$ where GLMY reports 0 (a triangle filled by an apex) and vice versa. On a synthetic lock-acquisition state cycle (`examples/state_machine/lock_cycle.json`) the directed analysis extracts the concrete generator $[s_0 \rightarrow s_1, s_1 \rightarrow s_2, s_2 \rightarrow s_0]$ with $\beta_1 = 1$ — a finding with no taint analogue.

13.3 Trust topologies, IAM and lateral movement

In cloud environments (AWS/Azure/Kubernetes) the worst breaches come from privilege-escalation *chains* and lateral movement (a compromised role in a frontend service assuming a secondary role with access to a secret bucket). The modeled object is the *identity-and-permission graph* (users, IAM roles, trust policies, services, network resources). Here the Fiedler cut (λ_2, f_2) optimally partitions the trust domains (e.g. DMZ vs. core); edges that cross the cut without intermediate controls surface as spectral anomalies at the trust boundary, and Forman–Ricci curvature — far from a detection error — *isolates the critical privilege bridges*: a strongly negative κ_F node is a role or credential that is the sole bridge between two disparate trust subnets (the ideal attacker pivot).

The agent then hands the suggested paths to Z3, which resolves the IAM-policy combinatorics to confirm whether the role-assumption chain is viable.

13.4 Security of autonomous agents and LLMs

Modern systems wiring LLMs to external tools (APIs, databases, shell, email) face *confused-deputy* and indirect prompt-injection attacks. The modeled object is the *tool-use graph*: user, model, context memory and external tools at different privilege levels, with edges typed `untrusted_input` (web data, read mail) and `privileged_execution` (drop a database, issue payments, run commands). The flaw is an authorization-invariant violation: an untrusted-input flow reaching a privileged-execution tool without an explicit sanitization or human-confirmation **gate**. This is exactly what the `llm-agent` adapter (`manifold.ingest.llm_agent`) already detects: a dangerous tool reachable from user input with no approval gate is a finding.

13.5 Mapping matrix and positioning

Formal / math layer	Modeled object	Mapped vulnerability
Auth invariant (Lean 4)	endpoints, decorators, API resources	BOLA / IDOR / BFLA
Directed homology (GLMY)	state-transition and async-call graph	reentrancy, state desync, deadlocks
Spectral (Fiedler)	network topologies, Kubernetes RBAC, IAM	trust-segmentation break, lateral movement
Discrete curvature (Ricci)	microservice and identity graphs	single points of failure, privilege bridges
Symbolic verifier (Z3)	path constraints in business logic	executable counterexamples (witnesses)

MANIFOLD is thus repositioned as a *formal engine for structural invariants and business logic*, not a competitor to Semgrep/CodeQL on local SQLi/XSS patterns. (1) It *abandons a sterile competition*: it does not try to beat taint on a 20-line servlet. (2) It *attacks what nobody solves*: design flaws, privilege coherence, concurrent protocols and agent architectures, where syntactic rules are silent because the code, line by line, looks legitimate. (3) It *keeps the software intact*: the typed IR, the Rust kernels, the Lean 4 verification, the Z3 orchestrator and the neuro-symbolic agent are all preserved, applied over the right relation graph. The narrative is therefore not “geometry lost to taint” but a single account of *which* program relations make structural invariants informative: a rigorously established negative boundary on linear injection SAST, and a positive one on the architectural flaws where taint is blind — both sides backed by working code and measured data.

References

- [1] M. Fiedler, *Algebraic connectivity of graphs*, 1973.
- [2] F. Chung, *Spectral Graph Theory*, 1997.
- [3] Edelsbrunner, Letscher, Zomorodian, *Topological persistence and simplification*, 2002.
- [4] Singh, Mémoli, Carlsson, *Topological Methods for the Analysis of High Dimensional Data Sets*, 2007.
- [5] Y. Ollivier, *Ricci curvature of Markov chains on metric spaces*, 2009.

- [6] P. Cousot & R. Cousot, *Abstract interpretation: a unified lattice model*, 1977.
- [7] S. Mac Lane, *Categories for the Working Mathematician*, 1971.
- [8] J. King, *Symbolic execution and program testing*, 1976.
- [9] E. M. Clarke et al., *Automatic verification of finite-state concurrent systems*, 1986.
- [10] F. Chung, *Laplacians and the Cheeger inequality for directed graphs*, 2005.
- [11] Grigor’yan, Lin, Muranov, Yau, *Homotopy theory for digraphs*, 2012.
- [12] M. Cuturi, *Sinkhorn distances: lightspeed computation of optimal transport*, 2013.
- [13] R. Forman, *Bochner’s method for cell complexes and combinatorial Ricci curvature*, 2003.
- [14] L. de Moura & N. Bjørner, *Z3: An efficient SMT solver*, 2008.
- [15] Y. Shoshitaishvili et al., *SoK: (State of) The Art of War: Offensive Techniques in Binary Analysis (angr)*, 2016.
- [16] OWASP Foundation, *OWASP Benchmark v1.2*, 2016–2024.
- [17] OWASP Foundation, *OWASP API Security Top 10* (API1:2023 Broken Object Level Authorization), 2023.
- [18] NIST, *Software Assurance Reference Dataset (SARD) / Juliet Test Suite*, 2017.
- [19] BerriAI, *LiteLLM: a unified interface to 100+ LLMs*, 2023–2024.
- [20] SonarSource, *SonarQube* (Community Edition, SonarJava security rules).
- [21] GitHub, *CodeQL* and the `java-security-extended` query suite.