

# MANIFOLD

## A Mathematical Engine for Mapping Vulnerabilities

MANIFOLD Project

Whitepaper v0.1 — working draft, not peer reviewed

### Abstract

MANIFOLD treats a software artifact not as a bag of rules to match, but as a *mathematical object*: a typed, weighted graph enriched with *algebraic*, *spectral*, *topological* and *geometric* structure. Over this object we compute a scalar field of **vulnerability potential**  $V(x)$ , which surfaces anomalous regions as geometric / topological / spectral *features* rather than as literal string matches. An **autonomous LLM agent** navigates the resulting “manifold”, formulates hypotheses about those regions, and dispatches **formal verifiers** (SMT, symbolic execution, abstract interpretation) that confirm or refute each hypothesis — closing the loop between statistical intuition and mathematical proof. This document defines the intermediate representation, formalizes each mathematical layer, states the mapping laws that connect mathematical features to vulnerability classes, and specifies the agent architecture and roadmap.

## 1 Motivation

Traditional scanners (SAST/DAST) enumerate syntactic patterns. They are brittle: they miss novel vulnerabilities and drown analysts in false positives. Two capabilities are missing:

- (1) A *unified, quantitative representation* of “program structure” onto which many *independent* signals can be projected and fused.
- (2) A *reasoning loop* that converts “suspicious region” into “proven vulnerability” — the LLM supplies the intuition, mathematics supplies the proof.

MANIFOLD provides both.

## 2 The core idea




Figure 1: The pipeline: an artifact becomes a typed graph, enriched with mathematical structure into a vulnerability scalar field  $V(x)$ , which an LLM agent navigates; a formal verifier confirms or refutes each hypothesis and the result re-embeds into the manifold.

A vulnerability is *not* a pattern; it is a *violation of a structural invariant* (a broken naturality condition, a persistent cycle, a negative curvature chokepoint, a taint flow across a spectral trust cut). Each layer detects a different *kind* of invariant violation, and the fused field  $V(x)$  ranks regions for the agent to inspect.

### 3 The Intermediate Representation (IR)

**Definition 1 (Program graph).** A *program graph* is a tuple  $G = (V, E, \tau_V, \tau_E, \omega)$  where

- $V$  is a finite set of *nodes* (modules, functions, blocks, statements, calls, variables, sources, sinks);
- $E \subseteq V \times V$  is a set of directed *edges*;
- $\tau_V : V \rightarrow \mathbb{N}_K$  assigns each node a *kind* from  $K = \{\text{module, function, block, statement, call, assign, variable, parameter, source, sink, gate}\}$ ;
- $\tau_E : E \rightarrow \mathcal{T}$  assigns each edge a *type*  $\mathcal{T} = \{\text{control, data, call, taint, trust, auth}\}$ ;
- $\omega : E \rightarrow \mathbb{R}_{\geq 0}$  assigns each edge a non-negative *weight* (default 1).

The IR is the single source of truth. Every adapter (source code, binary, web API, LLM agent) must produce a program graph; every analysis consumes it. The schema is language-agnostic and JSON-serializable (the Python and Rust layers share the identical schema).

## 4 Mathematical layers

For each layer we fix a subgraph by edge type  $\tau_E$ , then compute structure.

### 4.1 Spectral layer (algebraic graph theory)

For an edge type  $t$ , define the symmetric *adjacency*  $A_t$  and *degree*  $D_t$ , and the *combinatorial Laplacian*

$$L_t = D_t - A_t.$$

$L_t$  is positive semidefinite; its spectrum  $0 = \lambda_1 \leq \lambda_2 \leq \dots$  encodes global structure:

- $\lambda_2$  is the *algebraic connectivity* (Fiedler value). Its eigenvector  $f_2$  — the *Fiedler vector* — induces a canonical cut  $S^+ = \{v : f_2(v) \geq 0\}$ .
- The *spectral embedding*  $\Phi^{(d)}(v) = (f_2(v), f_3(v), \dots, f_{d+1}(v))$  maps nodes into  $\mathbb{R}^d$ .
- The Cheeger inequality bounds the bottleneck ratio  $h(G)$  by  $\frac{\lambda_2}{2} \leq h(G) \leq \sqrt{2 d_{\max} \lambda_2}$ .

**Signal 1 (spectral anomaly).** The reconstruction error of a node under a low-rank eigenmap reconstruction, and its distance from the Fiedler cut, are used as anomaly scores.

## 4.2 Topological layer (persistent homology)

From the program graph we build a *filtration* of simplicial complexes (e.g. a Vietoris–Rips or flag complex over the shortest-path metric, or a filtration driven by  $\omega$ ). Persistent homology yields, for each dimension  $k$ , a *persistence diagram*  $\text{Dgm}_k$  of birth–death pairs.  $\beta_0$  counts components,  $\beta_1$  counts independent cycles; points far from the diagonal have high persistence. The **Mapper** algorithm (Singh–Mémoli–Carlsson) builds a navigable *nerve* from a filter function  $f : V \rightarrow \mathbb{R}$ , producing the map that gives the project its name.

## 4.3 Geometric layer (discrete curvature)

The *Ollivier–Ricci curvature* of an edge  $(u, v)$  is

$$\kappa(u, v) = 1 - \frac{W_1(m_u, m_v)}{d(u, v)},$$

with lazy random-walk measures  $m_u, m_v$  and Wasserstein-1 distance  $W_1$ . Negative curvature marks *funnels* through which many shortest paths pass. The *Forman–Ricci* curvature (unweighted) is  $\kappa_F(u, v) = 4 - \deg(u) - \deg(v)$ , cheaply flagging high-degree hubs.

## 4.4 Algebraic layer (lattices & category theory)

*Abstract interpretation.* The *taint lattice*  $\mathbb{T} = \{\perp = \text{clean} \sqsubseteq \top = \text{tainted}\}$  (a product lattice over tags) gives a *Galois connection*  $(\alpha, \gamma)$  between the concrete collecting semantics and the abstract taint domain, guaranteeing soundness. *Category theory.* The program is a category **Prog**; effects are monads, authorization is a functor  $F : \mathbf{Prog} \rightarrow \mathbf{Auth}$ . **Design law (naturality of taint):** a vulnerability is a taint flow that breaks the naturality square of  $F$ .

## 4.5 Formal layer (symbolic execution & model checking)

Symbolic execution computes a path condition  $\Phi$  and symbolic state  $\sigma$ ; an SMT solver checks  $\text{SAT}(\Phi \wedge \text{bad})$ , whose model is a concrete exploit. Reachability is the CTL formula **EF** bad; taint propagation is a finite automaton whose accepted language is exactly the set of source  $\rightarrow$  sink paths.

# 5 The vulnerability scalar field

$$V(x) = \sigma\left(\sum_i \alpha_i s_i(x)\right),$$

where  $s_i$  are normalized layer signals (spectral anomaly, persistence outlier, curvature, lattice height, formal reachability),  $\alpha_i$  are (learnable) weights, and  $\sigma$  a logistic squeeze.  $V$  is a ranking, a Mapper filter, and a cost field for navigation.




Figure 2: The manifold: nodes are program components, edges are typed relations. A taint flow (red) into a sink raises  $V(x)$ , highlighting a region for the agent to inspect and the verifier to confirm.

## 6 Mapping laws

| #  | Law                                                   | Vulnerability class               |
|----|-------------------------------------------------------|-----------------------------------|
| L1 | Chokepoint ( $\kappa \ll 0$ )                         | privilege escalation, auth bypass |
| L2 | Cycle (persistent $H_1$ )                             | reentrancy, infinite recursion    |
| L3 | Trust cut (Fiedler cut)                               | injection across trust boundary   |
| L4 | Naturality break                                      | arbitrary taint violation         |
| L5 | Persistence outlier                                   | real vs. spurious features        |
| L6 | Reachability ( $\text{SAT}(\Phi \wedge \text{bad})$ ) | concrete exploit                  |

## 7 The autonomous LLM agent

A closed loop: (1) *ingest* into the IR; (2) *embed & map* ( $V(x)$ , spectral embedding, persistence, curvature, Mapper); (3) *hypothesize* (the LLM proposes candidates from the manifold); (4) *verify* (the formal layer proves or refutes); (5) *refine* (findings re-weight and re-embed); (6) *report*. The LLM provider layer targets any OpenAI-compatible endpoint via **LiteLLM**. The LLM does not do raw detection: mathematics proposes regions, the LLM proposes hypotheses, the formal layer proposes proofs.

## 8 Design challenges and mitigations

A design review identified five structural risks. For each we state the mitigation adopted in the prototype and the open problem.

**C1 – The bottleneck paradox.** Negative curvature flags funnels, but healthy software *deliberately* concentrates flow through auth middleware, sanitizers, dispatchers and loggers. Reporting these as privilege-escalation leads yields systematic false positives. *Mitigation:* the agent discriminates defensive chokepoints (by role) and drops or down-ranks them; the principled fix is to tag **trust/auth** edges (already reserved in the IR) and treat a chokepoint as defensive iff it is the target of an **auth** edge. *Open:* learning the defensive/offensive label from data.

**C2 – Symmetrization destroys causality.** The combinatorial Laplacian  $L = D - A$  and standard persistent homology act on undirected metric spaces, but control and data flow are directed:  $A$  tainting  $B$  is not  $B$  tainting  $A$ . Symmetrizing can destroy the direction of taint. *Mitigation:* the prototype implements Chung’s directed Laplacian (largest strongly connected component,

Perron vector, directed algebraic connectivity), which is undefined for a DAG — itself a meaningful signal. *Open*: directed flag complexes / path homology (Grigor’yan–Lin–Muranov–Yau).

**C3 – Cost of optimal transport.** Ollivier–Ricci requires a transportation LP per edge:  $\mathcal{O}(|E| \cdot \text{LP})$ , prohibitive beyond  $\sim 10^5$  edges. *Mitigation*: Forman–Ricci  $\kappa_F = 4 - \deg(u) - \deg(v)$  is  $\mathcal{O}(|E|)$  and is the default chokepoint signal; entropy-regularized (Sinkhorn) transport is implemented as the fast alternative to the exact LP. *Open*: adaptive regularization schedules for very large graphs.

**C4 – Operationalizing L4 (naturality).** A constructive rule is required. Let **auth** edges  $E_A$  mark privilege boundaries. Define the authorization functor  $F$  to label a node by whether it is reachable *behind* an auth edge. Then a taint flow  $s \rightarrow^* k$  *violates naturality* iff its path crosses at least one auth edge — i.e. untrusted data enters a protected region without the boundary neutralizing it. This is a reachability query over the combined  $(\text{data} \cup \text{taint} \cup \text{auth})$  graph, directly computable and automatable. *Implemented*: an authorization extractor emits **gate** nodes and **auth** edges from privilege-gate decorators, and a taint finding inside a gated function is reported as a naturality violation.

**C5 – Manifold-to-verifier bridge.** Geometric features are quantitative; Z3 needs exact preconditions. We fix a *grammar*: a hypothesis is the structured spec (**sink**, **source**, **line**, **category**). The manifold supplies the spec (it knows the sink and line), the LLM supplies the interpretation (CWE, description, confidence), and the verifier consumes the spec — never the LLM prose. This keeps the SMT query precise and reproducible.

The underlying thesis stands: static analyzers fail because they match words instead of invariants in a topological structure. The criterion is whether  $V(x)$  prunes the search space for Z3/angr more effectively than a plain DFS over the call graph.

## 9 Implementation status

| Phase | Deliverable                             | Status  |
|-------|-----------------------------------------|---------|
| F0    | Whitepaper, architecture                | done    |
| F1    | IR + Python SAST ingest + taint         | done    |
| F2    | Spectral + geometric kernels (Rust)     | done    |
| F3    | TDA (persistent homology, Mapper)       | done    |
| F4    | Formal verification (taint lattice, Z3) | done    |
| F5    | Autonomous LLM agent (LiteLLM)          | done    |
| F6    | Visualization + end-to-end demo         | done    |
| F7    | Binary / web / LLM adapters             | pending |

## 10 Validation and responsible use

Validation uses toy micro-benchmarks with known ground truth, then labeled corpora (OWASP Benchmark, SARD) once F4 lands, measuring precision/recall of the fused field versus raw taint. MANIFOLD is a defensive / authorized-testing tool; we release under the assumption of authorized use and encourage coordinated disclosure.

## References

- [1] M. Fiedler, *Algebraic connectivity of graphs*, 1973.

- [2] F. Chung, *Spectral Graph Theory*, 1997.
- [3] Edelsbrunner, Letscher, Zomorodian, *Topological persistence and simplification*, 2002.
- [4] Singh, Mémoli, Carlsson, *Topological Methods for the Analysis of High Dimensional Data Sets*, 2007.
- [5] Y. Ollivier, *Ricci curvature of Markov chains on metric spaces*, 2009.
- [6] P. Cousot & R. Cousot, *Abstract interpretation: a unified lattice model*, 1977.
- [7] S. Mac Lane, *Categories for the Working Mathematician*, 1971.
- [8] J. King, *Symbolic execution and program testing*, 1976.
- [9] E. M. Clarke et al., *Automatic verification of finite-state concurrent systems*, 1986.
- [10] F. Chung, *Laplacians and the Cheeger inequality for directed graphs*, 2005.
- [11] Grigor'yan, Lin, Muranov, Yau, *Homotopy theory for digraphs*, 2012.
- [12] M. Cuturi, *Sinkhorn distances: lightspeed computation of optimal transport*, 2013.